

Jingyu (Jack) Zhang

Department of Computer Science · Johns Hopkins University · Baltimore, MD

🌐 jackz.io

✉ jzhan237@jhu.edu

📞 443-453-8579

🔗 jackjyzhang

RESEARCH INTERESTS

My research focuses on the post-training and alignment of foundation models and agents, developing methods that simultaneously advance model capability alongside safety and adversarial robustness. My recent work leverages reinforcement learning to enable [multi-agent collaboration](#) and [adaptation to diverse values](#). My long-term goal is to build safe, collaborative agentic systems that reliably accomplish long-horizon, economically valuable tasks.

EDUCATION

Johns Hopkins University

Aug 2023 - Present

Ph.D. in Computer Science

Advisors: Daniel Khashabi and Benjamin Van Durme

Johns Hopkins University

Aug 2019 - May 2023

B.S. in Computer Science

GPA: 3.97/4.00

Additional Majors: Mathematics; Applied Mathematics & Statistics. Minor: Economics

PUBLICATIONS & PREPRINTS

Selected Works.....

- The Alignment Waltz: Jointly Training Agents to Collaborate for Safety**
Jingyu Zhang, Haozhu Wang, Eric Michael Smith, Sid Wang, Amr Sharaf, Mahesh Pasupuleti, Benjamin Van Durme, Daniel Khashabi, Jason Weston, Hongyuan Zhan.
ICLR 2026 submission. Preprint at <https://arxiv.org/abs/2510.08240>
- Controllable Safety Alignment: Inference-Time Adaptation to Diverse Safety Requirements**
Jingyu Zhang, Ahmed Elgohary, Ahmed Magooda, Daniel Khashabi, Benjamin Van Durme.
ICLR 2025
- Verifiable by Design: Aligning Language Models to Quote from Pre-Training Data**
Jingyu Zhang, Marc Marone, Tianjian Li, Benjamin Van Durme, Daniel Khashabi.
NAACL 2025 (oral presentation).
- SemStamp: A Semantic Watermark with Paraphrastic Robustness for Text Generation**
Abe Bohan Hou*, **Jingyu Zhang*** Tianxing He*, Yichen Wang, Yung-Sung Chuang, Hongwei Wang, Lingfeng Shen, Benjamin Van Durme, Daniel Khashabi, Yulia Tsvetkov.
NAACL 2024. [*Equal Contribution]

All Papers.....

- Beyond Reasoning Gains: Mitigating General Capabilities Forgetting in Large Reasoning Models**
Hoang Phan, Xianjun Yang, Kevin Yao, **Jingyu Zhang**, Shengjie Bi, Xiaocheng Tang, Madian Khabsa, Lijuan Liu, Deren Lei.
ICLR 2026 submission. Preprint at <https://arxiv.org/abs/2510.21978>.
- Jailbreak Distillation: Renewable Safety Benchmarking**
Jingyu Zhang, Ahmed Elgohary, Xiawei Wang, A S M Iftexhar, Ahmed Magooda, Benjamin Van Durme, Daniel Khashabi, Kyle Jackson.
Findings of EMNLP 2025
- Certified Mitigation of Worst-Case LLM Copyright Infringement**
Jingyu Zhang, Jiacan Yu, Marc Marone, Benjamin Van Durme, Daniel Khashabi.
EMNLP 2025
- Can A Society of Generative Agents Simulate Human Behavior and Inform Public Health Policy? A Case Study on Vaccine Hesitancy**
Abe Bohan Hou, Hongru Du, Yichen Wang, **Jingyu Zhang**, Zixiao Wang, Paul Pu Liang, Daniel Khashabi, Lauren Gardner, Tianxing He.
COLM 2025.
- Rationalyst: Pre-training Process-Supervision for Improving Reasoning**

- Dongwei Jiang, Guoxuan Wang, Yining Lu, Andrew Wang, **Jingyu Zhang**, Chuyu Liu, Benjamin Van Durme, Daniel Khashabi.
ACL 2025
10. **Core: Robust Factual Precision Scoring with Informative Sub-Claim Identification**
Zhengping Jiang, **Jingyu Zhang**, Nathaniel Weir, Seth Ebner, Miriam Wanner, Kate Sanders, Daniel Khashabi, Anqi Liu and Benjamin Van Durme.
Findings of ACL 2025
 11. **Self-(In)Correct: LLMs Struggle with Refining Self-Generated Responses.**
Dongwei Jiang, **Jingyu Zhang**, Orion Weller, Nathaniel Weir, Benjamin Van Durme, Daniel Khashabi.
AAAI 2025
 12. **TurkingBench: A Challenge Benchmark for Web Agents**
Kevin Xu, Yeganeh Kordi, Kate Sanders, Yizhong Wang, Adam Byerly, **Jingyu Zhang**, Benjamin Van Durme, Daniel Khashabi.
NAACL 2025
 13. **DiffNorm: Self-Supervised Normalization for Non-autoregressive Speech-to-speech Translation.**
Weiting Tan, **Jingyu Zhang**, Lingfeng Shen, Daniel Khashabi, Philipp Koehn.
NeurIPS 2024.
 14. **k-SemStamp: A Clustering-Based Semantic Watermark for Detection of Machine-Generated Text**
Abe Bohan Hou, **Jingyu Zhang**, Yichen Wang, Daniel Khashabi, Tianxing He.
Findings of ACL 2024.
 15. **The Language Barrier: Dissecting Safety Challenges of LLMs in Multilingual Contexts**
Lingfeng Shen, Weiting Tan, Sihao Chen, Yunmo Chen, **Jingyu Zhang**, Haoran Xu, Boyuan Zheng, Philipp Koehn, Daniel Khashabi.
Findings of ACL 2024.
 16. **On the Zero-Shot Generalization of Machine-Generated Text Detectors**
Xiao Pu, **Jingyu Zhang**, Xiaochuang Han, Yulia Tsvetkov, Tianxing He.
Findings of EMNLP 2023.
 17. **On the Blind Spots of Model-Based Evaluation Metrics for Text Generation**
Tianxing He*, **Jingyu Zhang***, Tianle Wang, Sachin Kumar, Kyunghyun Cho, James Glass, Yulia Tsvetkov.
ACL 2023 (oral presentation). [*Equal Contribution]
 18. **Geo-Seq2seq: Twitter User Geolocation on Noisy Data through Sequence to Sequence Learning**
Jingyu Zhang, Alexandra DeLucia, Chenyu Zhang, Mark Dredze.
Findings of ACL 2023.
 19. **PCFG-based Natural Language Interface Improves Generalization for Controlled Text Generation**
Jingyu Zhang, James Glass, Tianxing He.
*SEM 2023. Preliminary version accepted at 2nd Workshop on Efficient Natural Language and Speech Processing (ENLSP), NeurIPS 2022. **Best Paper Award.**
 20. **Changes in Tweet Geolocation over Time: A Study with Carmen 2.0**
Jingyu Zhang, Alexandra DeLucia, Mark Dredze.
8th Workshop on Noisy User-generated Text (W-NUT), COLING 2022.
 21. **Study of Manifestation of Civil Unrest on Twitter**
Abhinav Chinta*, **Jingyu Zhang***, Alexandra DeLucia, Anna L. Buzcak, Mark Dredze.
7th Workshop on Noisy User-generated Text (W-NUT), EMNLP 2021. [*Equal Contribution]

RESEARCH EXPERIENCE

Meta Superintelligence Labs - Trust Alignment Team Research Scientist Intern Mentors: Hongyuan Zhan and Jason Weston	May 2025 - Present (P/T after Aug 2025)
Microsoft Responsible AI Research Research Intern Mentor: Ahmed Elgohary Ghoneim	May 2024 - Apr 2025 (P/T after Aug 2024)
Center for Language and Speech Processing at Johns Hopkins University PhD Researcher Advisors: Daniel Khashabi and Benjamin Van Durme	Fall 2023 - Present
University of Washington Research Intern Advisors: Yulia Tsvetkov and Tianxing He	Summer 2022 - Spring 2023

MIT Computer Science and Artificial Intelligence Laboratory

Research Intern

Advisors: James Glass and Tianxing He

Spring 2022 - Summer 2022

Center for Language and Speech Processing at Johns Hopkins University

Undergraduate Researcher

Advisors: Mark Dredze and Benjamin Van Durme

Spring 2021 - May 2023

AWARDS & HONORS

- **Amazon AI Fellowship** - \$100000 stipend + tuition/fees over two years. Press coverage:    
- **Michael J. Muuss Research Award** - Top 1 out of 744 JHU CS undergraduates (2023). Press coverage: 
- **Best Paper Award** - ENLSP Workshop at NeurIPS 2022 (\$1000 prize).
- **CRA Outstanding Undergraduate Researcher Award Nominee** - 4 out of 744 JHU CS undergraduates (2022)
- **Pistrutto Research Fellowship** - \$4000 grant (Fall 2022). Press coverage: 
- **Bloomberg Distinguished Professor (BDP) Summer Program Recipient** - \$6000 grant (Summer 2021).
- **Upsilon Pi Epsilon** - International Honor Society for the Computing and Information Disciplines.
- **National Olympiad in Informatics in Provinces (NOIP)** - National 1st Prize Certification (2018)

INVITED TALKS

The Alignment Waltz: Jointly Training Agents to Collaborate for Safety

Invited talk, CLSP Seminar, Johns Hopkins University, Baltimore, Maryland

Oct 2025

Invited talk, OSU NLP Group, Ohio State University, Columbus, Ohio

Nov 2025

Jailbreak Distillation: Renewable Safety Benchmarking

Invited talk, Microsoft Responsible AI Tech Talks, Redmond, Washington

May 2025

Poster presentation, EMNLP 2025, Suzhou, China

Nov 2025

Controllable Safety Alignment: Inference-Time Adaptation to Diverse Safety Requirements

Invited talk, CLSP Seminar, Johns Hopkins University, Baltimore, Maryland

Oct 2024

Invited talk, Microsoft Responsible AI Tech Talks, Redmond, Washington

Dec 2024

Invited talk, ZhiDX AI Talk, Online

June 2025

Verifiable by Design: Aligning Language Models to Quote from Pre-Training Data

Oral presentation, the 11th Mid-Atlantic Student Colloquium on Speech, Language and Learning

May 2024

Invited talk, Microsoft Hallucination Deep Dive, Redmond, Washington

Aug 2024

Invited talk, Institute for Assured Autonomy, Johns Hopkins University, Baltimore, Maryland

Oct 2024

Oral presentation, NAACL 2025, Albuquerque, New Mexico

May 2025

SemStamp: A Semantic Watermark with Paraphrastic Robustness for Text Generation

Invited talk, CLSP Seminar, Johns Hopkins University, Baltimore, Maryland

Dec 2023

On the Blind Spots of Model-Based Evaluation Metrics for Text Generation

Oral Presentation, ACL 2023, Toronto, Canada

July 2023

PCFG-based Natural Language Interface Improves Generalization for Controlled Text Generation

Spotlight Presentation (Best Paper Award), ENLSP Workshop, NeurIPS 2022, New Orleans, USA

Dec 2022

INDUSTRY EXPERIENCE

Meta Platforms, Inc.

Meta Superintelligence Labs - Trust Alignment

Research Scientist Intern

May 2025 - Present (P/T after Aug 2025)

- Conducted research related to multi-agent safety alignment (ICLR 2026 submission: <https://arxiv.org/abs/2510.08240>).
- Implemented a multi-agent reinforcement learning framework to support model co-evolution.
- Collaborating with FAIR on enhancing the parallel reasoning capabilities for LLMs.

Microsoft Corporation

Responsible AI Research

Research Intern

May 2024 - Apr 2025 (P/T after Aug 2024)

- Conducted research related to safety alignment and automatic red teaming.
- Resulted in two papers: Controllable Safety Alignment (ICLR 2025, <https://arxiv.org/abs/2410.08968>), Jailbreak Distillation (Findings of EMNLP 2025, <https://arxiv.org/abs/2505.22037>).

ByteDance Ltd.

Lark Explorer Department

C++ Development Intern

May 2020 - May 2021 (P/T after Aug 2020)

- Carried out C++ cross-platform development interacting with macOS kernel and Windows Win32 API.
- Performed client-side development with Electron and Node.js related to performance optimization.
- Conducted data science analytics on extensive user-generated data with Apache Hive and Python.

TEACHING & MENTORSHIP EXPERIENCE

EN.601.465/665 Natural Language Processing

Course Assistant

JHU Department of Computer Science

Fall 2021, Fall 2022

- EN.601.465/665 is a mixed graduate / upper-level undergraduate course in NLP taught by Prof. Jason Eisner.
- Conducted grading of homework and exam papers and held review sessions and office hours on a weekly basis.
- Received an average score of 4.80/5.00 on student TA evaluation (100% “Good” or “Excellent” rating).

Code in Place 2021

Section Leader (Volunteer)

Stanford University Department of Computer Science

April 2021 - May 2021

- Worked with a team of more than 50 teaching leads and 1000 section leaders to support 10,000+ students across the world as they navigate the first five weeks of CS 106A: Programming Methodology course.
- Prepared materials and taught a Python programming section of 10 students on a weekly basis.

Students Mentored

- Abe Bohan Hou [✉](#), undergraduate student at JHU. Now PhD student at Stanford University.
- Matt Hexuan Wang, undergraduate student at JHU.
- Jiacan Yu, master’s student at JHU.

ACADEMIC SERVICE

- Application Mentor, JHU CLSP pre-application support program (2023-present)
- Curriculum Committee, Department of Computer Science, Johns Hopkins University (2023-present)
- Recruitment Committee, Center for Language and Speech Processing, Johns Hopkins University (2023-present)

SKILLS & COURSEWORK

Technical Skills

- Programming Languages: Python, C/C++, Java, OCaml, HTML, CSS, Javascript, MATLAB, SQL
- Frameworks: Huggingface, PyTorch, OpenRLHF, Sklearn, Pandas, Numpy, Electron
- Development Workflow: $\LaTeX$, Bash, Emacs, Git, Makefile, GN build

Related Coursework

- NLP/ML: Natural Language Processing, Machine Translation, Multilingual NLP, Artificial Agents, Machine Learning, Deep Learning, Human-Computer Interaction, Probabilistic Models of the Visual Cortex
- Math/Stats: Real Analysis, Abstract Algebra, Topology, Differential Equations, Probability, Statistics, Optimization, Time Series Analysis, Game Theory, Mathematical & Computational Foundations of Data Science