

Jingyu (Jack) Zhang

Department of Computer Science · Johns Hopkins University · Baltimore, MD

🌐 jackz.io

✉ jzhan237@jhu.edu

📞 443-453-8579

🐙 jackjyzhang

EDUCATION

Johns Hopkins University

2023 - Present

Ph.D. in Computer Science

Advisors: Daniel Khashabi and Benjamin Van Durme

Johns Hopkins University

2019 - 2023

B.S. in Computer Science

GPA: 3.97/4.00

Additional Majors: Mathematics; Applied Mathematics & Statistics. Minor: Economics

RESEARCH INTERESTS

AI safety and alignment, robust evaluation and scalable oversight, controllability of language models

PUBLICATIONS & PREPRINTS

- SemStamp: A Semantic Watermark with Paraphrastic Robustness for Text Generation**
Abe Bohan Hou*, **Jingyu Zhang*** Tianxing He*, Yichen Wang, Yung-Sung Chuang, Hongwei Wang, Lingfeng Shen, Benjamin Van Durme, Daniel Khashabi, Yulia Tsvetkov.
arXiv preprint. [*Equal Contribution]
- On the Zero-Shot Generalization of Machine-Generated Text Detectors**
Xiao Pu, **Jingyu Zhang**, Xiaochuang Han, Yulia Tsvetkov, Tianxing He.
In Proc. of Findings of EMNLP 2023.
- On the Blind Spots of Model-Based Evaluation Metrics for Text Generation**
Tianxing He*, **Jingyu Zhang***, Tianle Wang, Sachin Kumar, Kyunghyun Cho, James Glass, Yulia Tsvetkov.
In Proc. of ACL 2023. Oral Presentation. [*Equal Contribution]
- Geo-Seq2seq: Twitter User Geolocation on Noisy Data through Sequence to Sequence Learning**
Jingyu Zhang, Alexandra DeLucia, Chenyu Zhang, Mark Dredze.
In Proc. of Findings of ACL 2023.
- PCFG-based Natural Language Interface Improves Generalization for Controlled Text Generation**
Jingyu Zhang, James Glass, Tianxing He.
*In Proc. of *SEM 2023*. Preliminary version accepted at *2nd Workshop on Efficient Natural Language and Speech Processing (ENLSP)*, *NeurIPS 2022. Best Paper Award*.
- Changes in Tweet Geolocation over Time: A Study with Carmen 2.0**
Jingyu Zhang, Alexandra DeLucia, Mark Dredze.
In Proc. of the 8th Workshop on Noisy User-generated Text (W-NUT), *COLING 2022*.
- Study of Manifestation of Civil Unrest on Twitter**
Abhinav Chinta*, **Jingyu Zhang***, Alexandra DeLucia, Anna L. Buzcak, Mark Dredze.
In Proc. of the 7th Workshop on Noisy User-generated Text (W-NUT), *EMNLP 2021*. [*Equal Contribution]

AWARDS & HONORS

- Best Paper Award - ENLSP Workshop at NeurIPS 2022
- Michael J. Muuss Research Award - 1 out of 744 JHU CS undergraduates. Press coverage: [🔗](#)
- CRA Outstanding Undergraduate Researcher Award Nominee - 4 out of 744 JHU CS undergraduates
- Pistrutto Research Fellowship - \$4000 grant, Fall 2022. Press coverage: [🔗](#)
- Bloomberg Distinguished Professor (BDP) Summer Program Recipient - \$6000 grant, Summer 2021
- Upsilon Pi Epsilon - International Honor Society for the Computing and Information Disciplines
- National Olympiad in Informatics in Provinces (NOIP) - National 1st Prize Certification (2018)

RESEARCH EXPERIENCE

Center for Language and Speech Processing at Johns Hopkins University

Fall 2023 - Present

PhD Researcher

Advisors: Daniel Khashabi and Benjamin Van Durme

University of Washington

Summer 2022 - Spring 2023

Undergraduate Research Intern

Advisors: Yulia Tsvetkov and Tianxing He

MIT Computer Science and Artificial Intelligence Laboratory

Spring 2022 - Summer 2022

Undergraduate Research Intern

Advisors: James Glass and Tianxing He

Center for Language and Speech Processing at Johns Hopkins University

Spring 2021 - May 2023

Undergraduate Researcher

Advisors: Mark Dredze and Benjamin Van Durme

PRESENTATIONS

SemStamp: A Semantic Watermark with Paraphrastic Robustness for Text Generation

Lightning Talk, JHU CLSP Seminar, Baltimore, Maryland

Dec 2023

On the Blind Spots of Model-Based Evaluation Metrics for Text Generation

Oral Presentation, ACL 2023, Toronto, Canada

July 2023

PCFG-based Natural Language Interface Improves Generalization for Controlled Text Generation

Spotlight Presentation, ENLSP Workshop, NeurIPS 2022, New Orleans, USA

Dec 2022

ACADEMIC SERVICE

Reviewing

- Reviewer, ACL 2023 (Generation track)
- Reviewer, Workshop on Instruction Tuning and Instruction Following, NeurIPS 2023

Outreach

- Application Mentor, JHU CLSP pre-application support program (2023)
- Curriculum Committee, Department of Computer Science, Johns Hopkins University (2023)
- Recruitment Committee, Center for Language and Speech Processing, Johns Hopkins University (2023)

TEACHING EXPERIENCE

EN.601.465/665 Natural Language Processing

Course Assistant

JHU Department of Computer Science

Fall 2021, Fall 2022

- EN.601.465/665 is a mixed graduate / upper-level undergraduate course in NLP taught by Prof. Jason Eisner
- Conducted grading of homework and exam papers
- Held review sessions and office hours on a weekly basis
- Received an average score of 4.80/5.00 on student TA evaluation (100% “Good” or “Excellent” rating)

Code in Place 2021

Section Leader (Volunteer)

Stanford University Department of Computer Science

April 2021 - May 2021

- Worked with a team of more than 50 teaching leads and 1000 section leaders to support 10,000+ students across the world as they navigate the first five weeks of CS 106A: Programming Methodology course
- Prepared materials and taught a Python programming section of 10 students on a weekly basis

INTERNSHIP EXPERIENCE

ByteDance Ltd.

C++ Development Intern

Lark Explorer Department

May 2020 - May 2021 (P/T after Summer 2020)

- Carried out C++ cross-platform development interacting with macOS kernel and Windows Win32 API
- Performed client-side development with Electron and Node.js related to performance optimization
- Conducted data science analytics on extensive user-generated data with Apache Hive and Python

SKILLS & COURSEWORK

- **Programming Languages:** Python, C/C++, Java, OCaml, HTML, CSS, Javascript, MATLAB, SQL
- **Frameworks:** Huggingface, PyTorch, Sklearn, Pandas, Numpy, Electron, Windows/macOS native APIs
- **Development Workflow:** L^AT_EX, Bash, Emacs, Git, Makefile, GN build

Last Updated: December 2, 2023.